MACHINE LEARNING

patrick.schwabl@ifkw.lmu.de
@patrickschwabl
My LMU page

The tidymodels packages revisited

1. Train-test split

```r
# Split data into training and testing sets
set.seed(1234)
vote_split <- voters_select %>%
    rsample::initial_split(prop = 0.8,
                strata = turnout16_2016)

vote_train <- training(vote_split)
vote_test <- testing(vote_split)


dplyr::glimpse(vote_train)
dplyr::glimpse(vote_test)
```

2. Create recipe

```r
# Normalize predictors
vote_recipe <-
  recipes::recipe(turnout16_2016  ~ ., data = vote_train) %>%
  update_role(case_identifier, new_role = "ID") %>%
  themis::step_upsample(turnout16_2016) %>%
  recipes::step_normalize(all_numeric())


vote_recipe
```

3. Specify model

```r
## Specify a knn model
knn_spec <- parsnip::nearest_neighbor() %>%
    parsnip::set_engine("kknn") %>%
    parsnip::set_mode("classification")
```

4. Create workflow
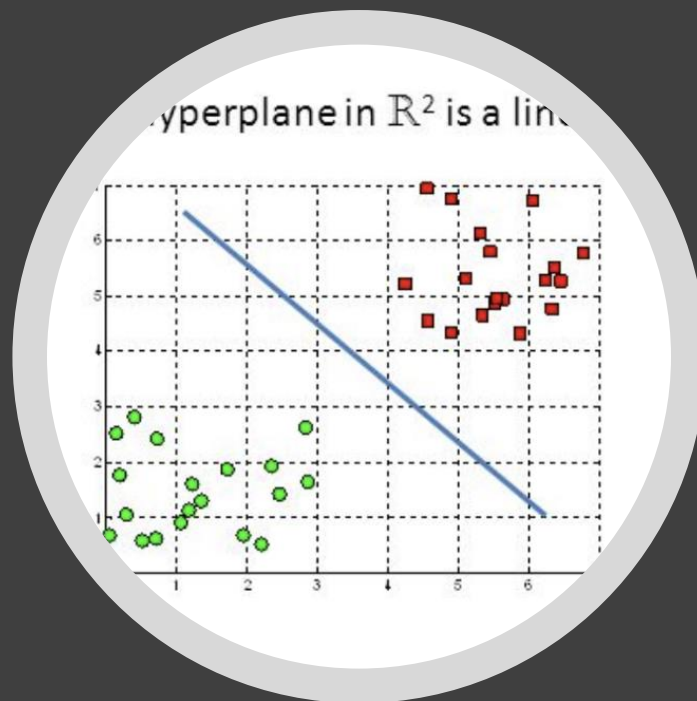
```r
## Add the recipe + model to a workflow
vote_workflow <- workflow() %>%
    workflows::add_recipe(vote_recipe) %>%
    workflows::add_model(knn_spec)


vote_workflow
```

5. Fit model

```r
# fit the final best model to the training set
 and evaluate the test set
knn_fit <- vote_workflow %>%
  tune::last_fit(vote_split)


print(knn_fit)
```
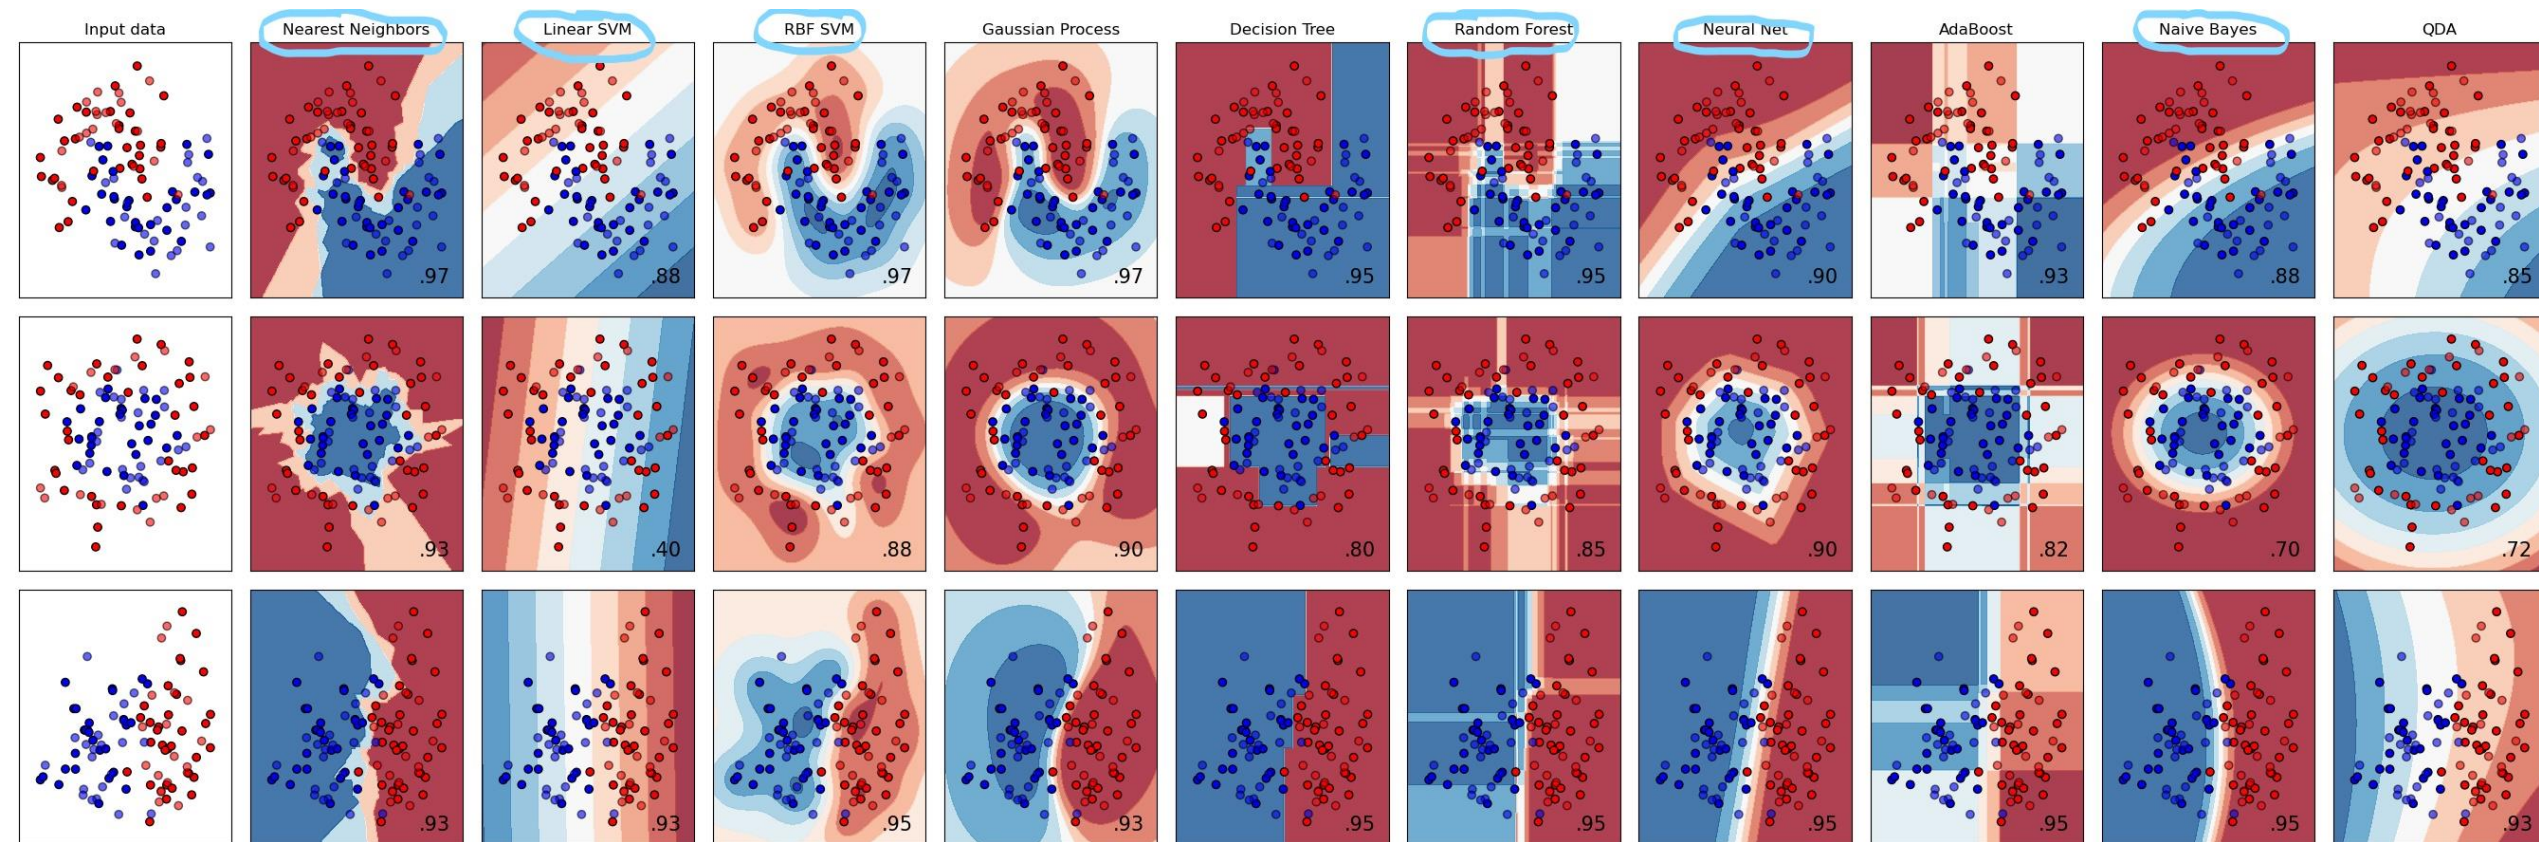
https://www.tidymodels.org/packages/
https://dplyr.tidyverse.org/

The zoo of supervised learning models

The zoo of supervised learning (classification) models

| Input data | Nearest Neighbors | Linear SVM | RBF SVM | Gaussian Process | Decision Tree | Random Forest | Neural Net | AdaBoost | Naive Bayes | QDA |

The zoo of supervised learning (classification) models – multinominal Naïve Bayes (NB)

Data:

Histogram of movie review data

Negative: "it's so laddish and juvenile, only teenage boys could possibly find it funny"

Positive: "provides a porthole into that noble, trembling incoherence that defines us all"

https://www.youtube.com/watch?v=O2L2Uv9pdDA
https://doi.org/10.3115/1219840.1219855
http://www.cs.cornell.edu/people/pabo/movie-review-data/
https://towardsdatascience.com/naive-bayes-classifier-81d512f50a7c
https://web.stanford.edu/~jurafsky/slp3/
https://towardsdatascience.com/k-nearest-neighbors-knn-explained-cbc31849a7e3

The algorithm:

1. Calculate a list (histogram) of relative frequencies (probabilities/likelihoods) per class using the *bag of words* approach.

2. Calculate probabilities of seeing a positive or negative review

$$p(positive) = \frac{5330}{10660} = 0.5$$

$$p(negative) = \frac{5330}{10660} = 0.5$$

3. Calculate conditional probabilities for a **single movie review**:

$$p(funny\ film|\ positive) = p(positive) * p(funny) * p(film) = 0.5 * 0.05 * 0.21 = 0.00525$$

$$p(funny\ film|\ negative) = p(negative) * p(funny) * p(film) = 0.5 * 0.00 * 0.17 = 0.00$$

→ Conditional probability scores of seeing the words "funny" and "film" given that the movie review is a positive or negative one
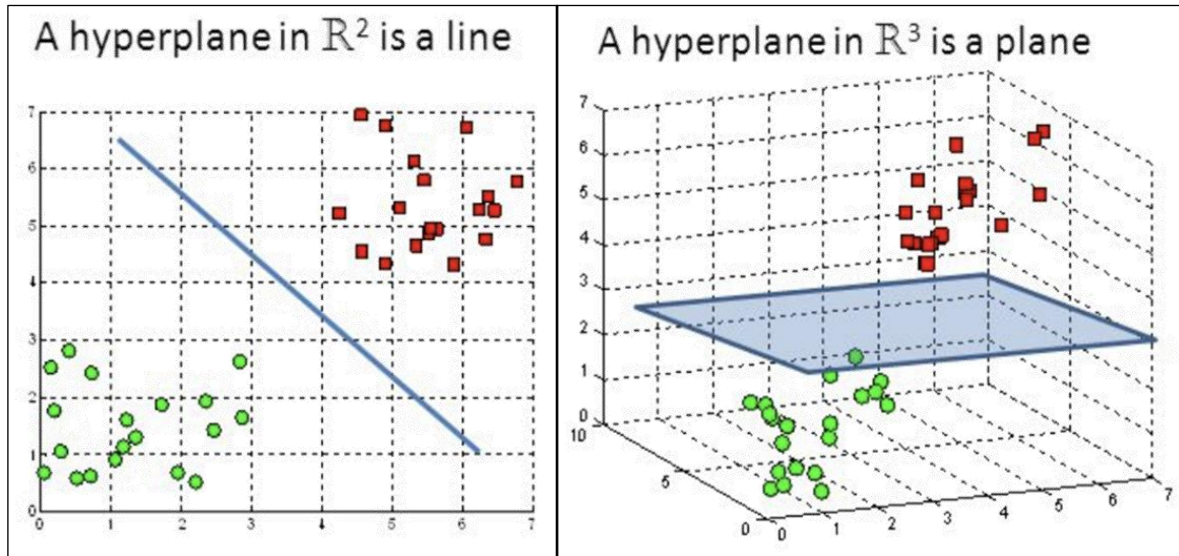
4. Compare scores for both classes: $0.00525 > 0.00$

5. Evaluate model on a testing dataset with known labels

The zoo of supervised learning (classification) models – Support Vector Machine (SVM)

Separating hyperplanes:

A hyperplane in $\mathbb{R}^2$ is a line

A hyperplane in $\mathbb{R}^3$ is a plane

Margins:

Small Margin

Large Margin

Support Vectors

https://towardsdatascience.com/support-vector-machine-introduction-to-machine-learning-algorithms-934a444fca47
https://medium.datadriveninvestor.com/support-vector-machines-svms-4bcccbd78369
http://vision.stanford.edu/teaching/cs231n-demos/linear-classify/
https://towardsdatascience.com/support-vector-machine-explained-8d75fe8738fd

The zoo of supervised learning (classification) models – Support Vector Machine (SVM)

The kernel trick:

Non-linear kernels:

https://towardsdatascience.com/svm-classifier-and-rbf-kernel-how-to-make-better-models-in-python-73bb4914af5b

The zoo of supervised learning (classification) models – Support Vector Machine (SVM)

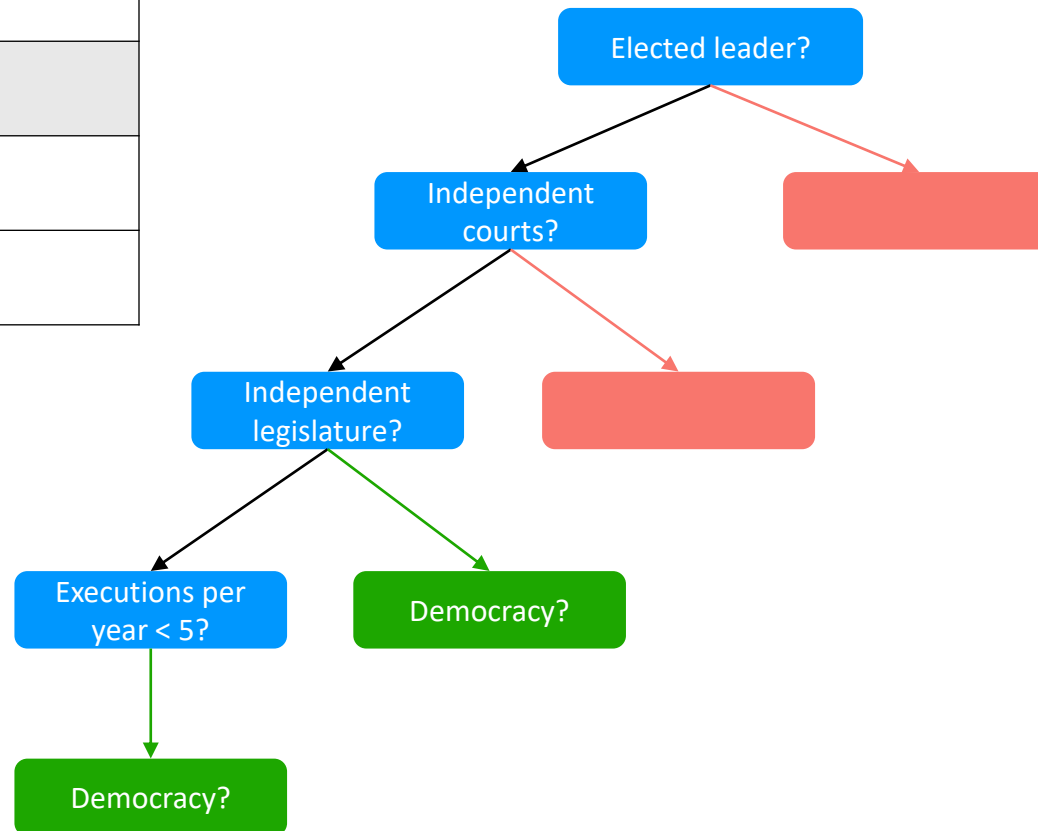**The algorithm for a linear SVM in two dimensions:**

1. "Draw" data points into a cartesian coordinate system with two dimensions

2. If the data are not linearly separable (you cannot draw a straight line that separates them) apply the kernel trick and "pull" the data apart into a third dimension

3. Find a straight line (or hyperplane) that separates the classes

4. Adjust the hyperplane until the average distance to the nearest points of both classes is maximized (maximum margin classifiers)

5. Evaluate model on a testing dataset with known labels

The zoo of supervised learning (classification) models – Random Forest

Original dataset

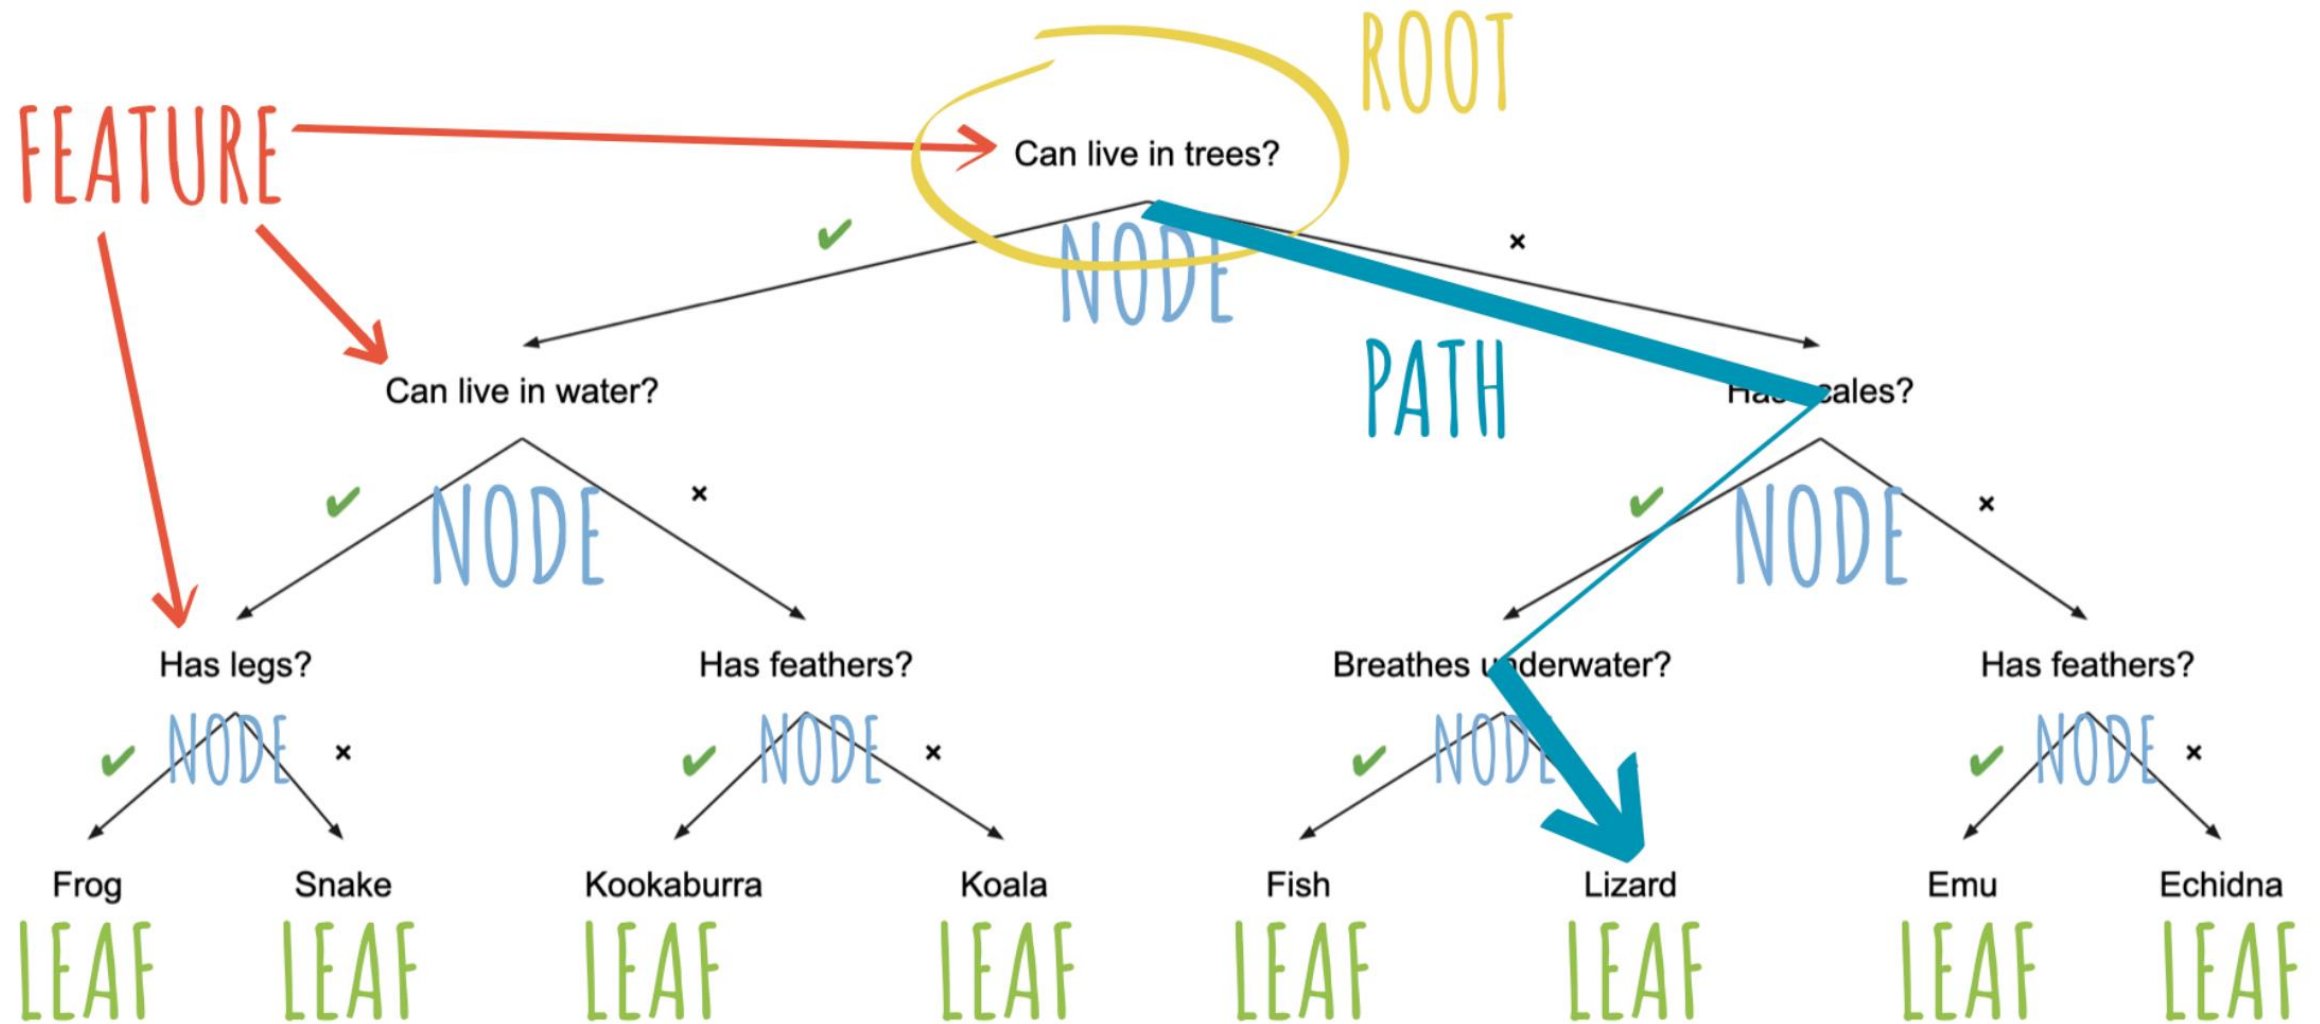| Elected leader	Independent courts	Independent legislature	Executions per year	Democracy
No	No	No	80	No
Yes	Yes	Yes	2	Yes
Yes	Yes	No	5	No
Yes	No	Yes	61	No

A decision tree

https://www.youtube.com/watch?v=J4Wdy0Wc_xQ
https://towardsdatascience.com/understanding-decision-trees-for-classification-python-9663d683c952
https://towardsdatascience.com/decision-tree-and-random-forest-explained-8d20ddabc9dd

The zoo of supervised learning (classification) models – Random Forest
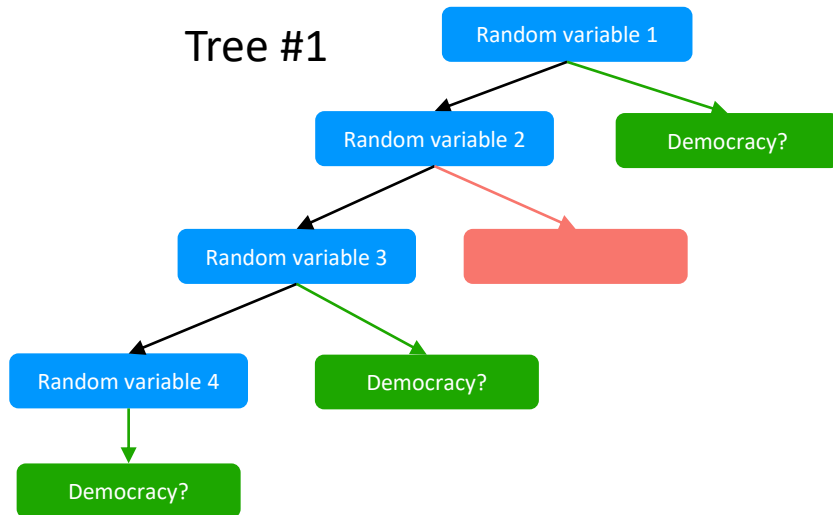
# The zoo of supervised learning (classification) models – Random Forest

New data point

| Elected leader	Independent courts	Independent legislature	Executions per year	Democracy
Yes	Yes	Yes	523	No

Tree #4

Random variable 1
Random variable 2
Democracy?
Random variable 3
Random variable 4
Democracy?
Democracy?

Tree #1

Random variable 1
Random variable 2
Democracy?
Random variable 3
Random variable 4
Democracy?
Democracy?

Tree #2

Random variable 1
Random variable 2
Random variable 3
Random variable 4
Democracy

Tree #3

Random variable 1
Random variable 2
Random variable 3
Random variable 4
Democracy?
Democracy?

Democracy?
Yes: 1
No: 3

The zoo of supervised learning (classification) models – Random Forest

The algorithm for a random forest with 10 trees

1. Create a bootstrapped dataset from the original dataset

2. Build a decision tree with the number of decision nodes equal to the number of predictors

3. Train the tree with the bootstrapped dataset

4. Repeat steps 1-3 10 times

5. Evaluate model on a testing dataset with known labels

    1. Let every tree decide on a new data point

    2. Let the trees vote on how to classify the new data point

The zoo of supervised learning models

Train test validate

V-fold cross validation

V-fold cross validation

Splitting the data for vfold validation

rsample::initial_split()

Sample rows are assigned at random

Training — For learning: to create and optimize the model

Testing — Only for performance evaluation **Not for Training!!!**

rsample:: **vfold_cv(train_data, v = 4)**

Sample rows are assigned at random

Cross Validation Folds – equal subsets of the training data

V-fold Validation

Cross Validation Folds

1 2 3 4

workflows::fit_resamples()

Fit model using different combinations of v-1 folds

1 2 3 4

1,2,3 | Test — Fit and Assess: Iteration 1

One fold is retained (blue) as a "Testing/Assessment set" to evaluate the performance of the model built on the other folds (yellow)

1 2 3 4

1,2,4 | Test — Fit and Assess: Iteration 2

1 2 3 4

1,3,4 | Test — Fit and Assess: Iteration 3

Performance **across iterations** is used to tuned parameters for optimal performance → **Final Model**

1 2 3 4

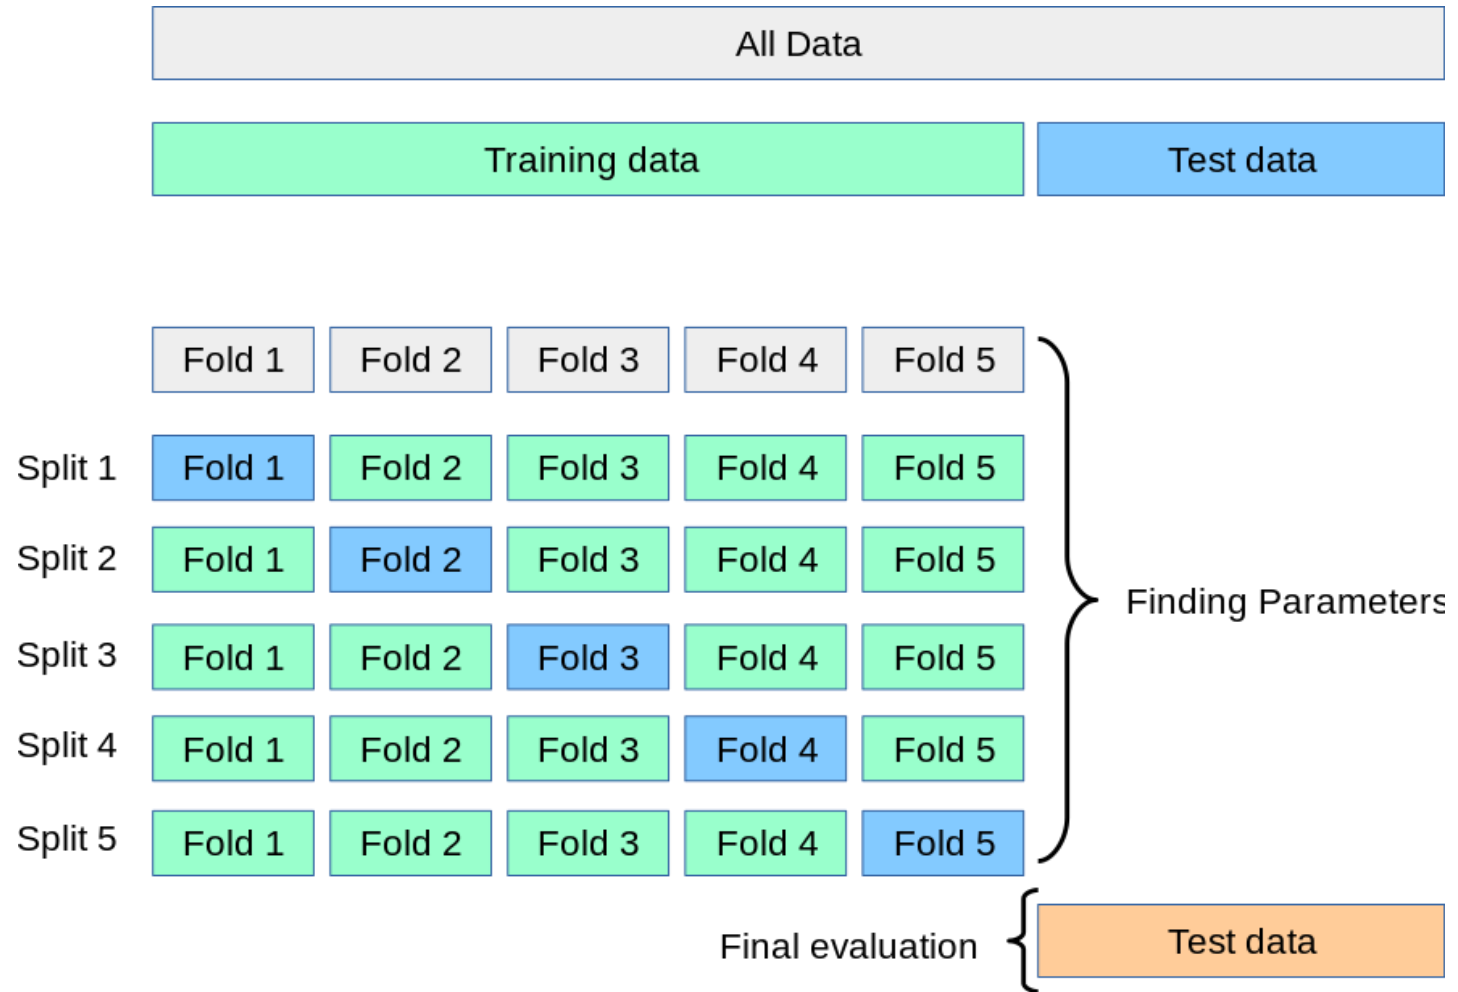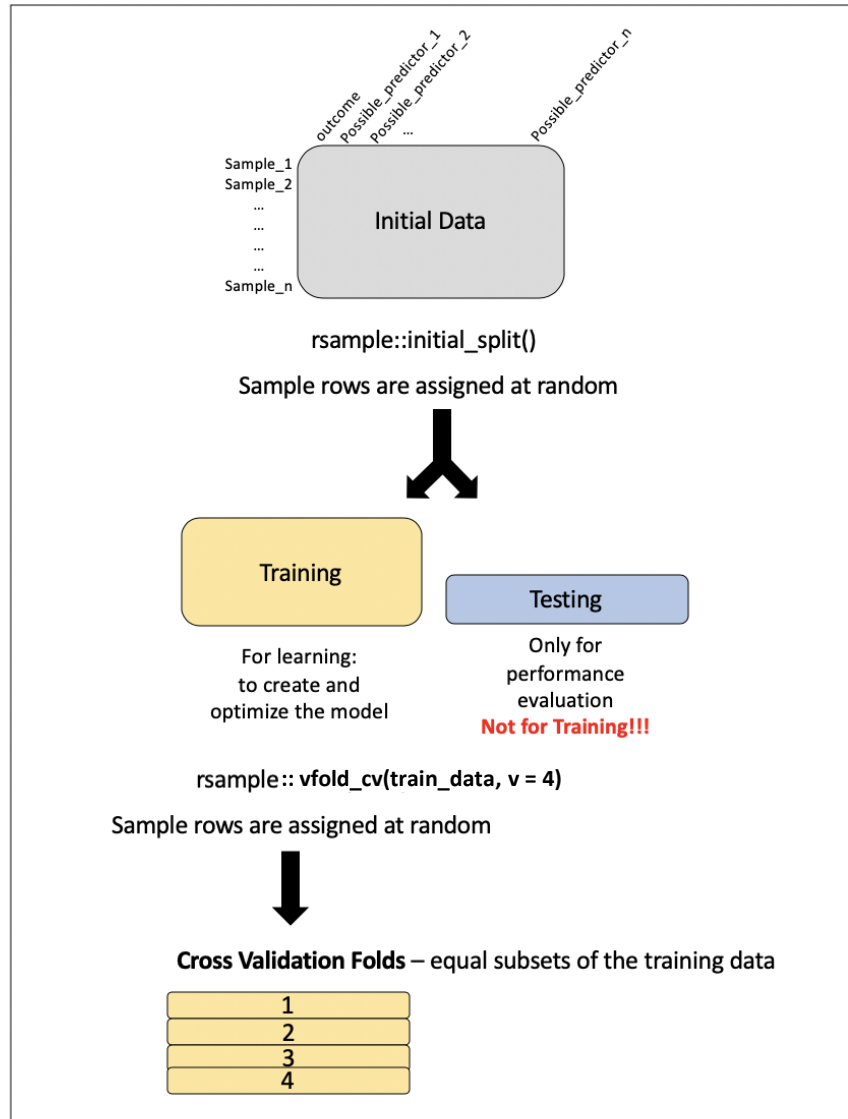2,3,4 | Test — Fit and Assess: Iteration 4

The initial testing set is used for the evaluation of the **final model** performance

Testing